

# CSE535: Asynchronous Systems

## Fall 2017

### Introduction

Scott D. Stoller



Stony Brook University

# Definition

- **Distributed system:** a collection of computers linked by a network and equipped with software that supports sharing of resources.
- “A **distributed system** is one in which the failure of a computer you didn’t even know existed can render your own computer unusable.”
  - Leslie Lamport, 2013 ACM Turing Award winner

# Motivation: Resource Sharing, Communication

## ● Resource Sharing

- ◆ Resources include hardware, software, and information
- ◆ Resource sharing increases utilization of resources hence is cost effective.
- ◆ Resource sharing allows higher total resources (e.g., RAM, CPU) than available in a single machine.

## ● Communication / Cooperation

- ◆ Examples: Skype, Facebook, SharePoint, medical information systems, factory automation systems, electronic stock exchange

# Motivation: Reliability

- Distributed systems may have **higher or lower** availability than centralized systems, depending on design.
- **Lower availability** arises when resource sharing introduces a single point of failure for multiple resources.
  - ◆ **Example:** In Cornell CS Department, all printers are networked. One day, the disk controller for the disk with the printer information database died. **No one could print**, from any computer to any printer, until it was replaced.
- **Higher availability** is generally achieved through replication (redundancy).

# Potential Disadvantages

- Lower speed, if network is slow
- Lower availability, if not designed for reliability
- Lower security (more points of vulnerability)
- Cost of network hardware and network administration
  - ◆ Savings from resource sharing usually outweighs this

# Design Goals: Transparency

- **Transparency:** the distributed system behaves like a (fast, reliable) centralized system.
  - ◆ The distribution is transparent (invisible) to users.
- Transparency makes distributed systems **easier to use**.
  - ◆ Centralized systems are simpler and more familiar.
- Transparency varies by software layer.
  - ◆ Higher layers (applications) generally have higher transparency.
  - ◆ Middle layers, i.e., middleware, provide transparency.

# Design Goals: Transparency

- Different aspects of behavior give rise to different aspects of transparency.
- **Location Transparency:** users do not need to know location of resources to access them.
  - ◆ Example: Print to any networked printer without knowing which print server it is connected to.
  - ◆ Example: Access my home directory without knowing which disk or file server it is stored on.
- **Host Transparency:** resources can be accessed the same way from all hosts.
  - ◆ `/home/facfs1/stoller` accesses my home directory on all Linux and Solaris machines in the CS Dept.
  - ◆ `/usr/bin` does not access the same directory on all machines. It shouldn't!

# Design Goals: Transparency

- **Concurrency Transparency:** concurrent access to a resource has the same outcome as sequential access, i.e., the same outcome as performing the operations sequentially in some total order (the “**serialization order**”).
- Exact meaning of concurrency transparency depends on
  - ◆ The set of operations, primarily their **granularity**
    - Smaller granularity provides weaker guarantees but typically allows more efficient implementation.
  - ◆ The requirements on acceptable serialization order
    - Looser requirements provide weaker guarantees (greater deviation from centralized behavior) but typically allow more efficient (e.g., lower latency) implementation.



# Design Goals: Transparency: Granularity

- **Example:** Processes P1 and P2 concurrently access a file F in a distributed filesystem (DFS).
- Assume sector size = 1KB and block size = 2KB.
- P1: write(0 1 2) || P2: write(A B C)
  - ◆ Each character above represents 1KB of data.
- What does sequential consistency imply about number of possible outcomes (POs)? Depends on what operations are considered. Sequential consistency for:
  - ◆ Sector operations (provided by disk controller):  $2^3$  POs
  - ◆ Block operations (provided by typical filesystem):  $2^2$  POs
  - ◆ Calls to write: 2 POs
- **Example:** NFS v3 provides block-level atomicity, with 4KB write blocks and 32 KB read blocks.

# Design Goals: Transparency: Order

- Typical requirements on serialization order, from strongest (most desirable) to weakest:
  - ◆ R1: Non-overlapping (in real time) operations must be serialized in real-time order.
  - ◆ R2: Same as above, except allow bounded violations, e.g., operations less than 1 sec apart can be serialized in either order.
  - ◆ R3: No requirements.
- **Example:** A file F in a DFS initially contains 0 1 2. Process P1 writes A B C, and then P2 reads from F. Can P2 read 0 1 2?
  - ◆ With R1, no. With R2, yes, if the interval between the operations is small enough. Example: NFS provides R2, because a process can read stale cached data.

# Design Goals: Reliability

- Common reliability **metrics**:
  - ◆ **Availability**: the fraction of requests that are processed in a timely manner
  - ◆ **MTTF**: Mean Time To Failure
- Reliability is typically achieved by **replication**.
- **Example**: File availability  $A_{\text{file}}$  in terms of disk availability.
  - ◆ **Without replication**:  $A1_{\text{file}} = A_{\text{disk}}$
  - ◆ **With replication of files on 2 disks**:
$$A2_{\text{file}} = 1 - \text{Prob}(\text{both disks down}) = 1 - (1 - A_{\text{disk}})^2$$
    - If  $A_{\text{disk}} = 0.99$ ,  $A1_{\text{file}} = 0.99$  and  $A2_{\text{file}} = 0.9999$ .

# Design Goals: Security, Scalability

## ● Security

- ◆ Includes **secrecy, integrity, availability**, ...
- ◆ Distributed systems often encompass systems in multiple administrative domains.
  - **Example:** medical information system involving hospital, doctor's office, insurance company, government agency (e.g., Medicare), etc.
- ◆ Policies for access control and delegation must express complex **trust relationships** between different **administrative domains**.

## ● Scalability

- ◆ Ability of the system to support more concurrent users or more tasks as more hardware resources are added.