

CSE535: Asynchronous Systems

Fall 2017

Distributed System Models

Scott D. Stoller



Stony Brook University

What are (Distributed System) Models?

- A **model** characterizes a class of systems.
- A **good model** includes all **relevant** aspects and ignores **irrelevant** details.
- Different models for different purposes, e.g., functional model, performance model, real-time model
- **Sequential systems:** less variety in the models
 - ◆ Theoretical models: Turing machine, lambda calculus
 - ◆ More detailed models: CPU, bus, RAM, disk
- **Distributed systems:** wider variety of models
- **Dimensions:**
 - ◆ Network Structure
 - ◆ Synchrony (Timing)
 - ◆ Failures

Network Structure


- **Network structure:** Pattern of interconnection of hosts
- Often represented by a graph.
- May be static or dynamic (e.g., reconfigurable).
- May be initially known to some or all hosts.
- May be connected or partitioned.
- **Examples:**
 - ◆ Bus, e.g., Ethernet
 - ◆ Ring, e.g., token ring, FDDI
 - ◆ Switched network, e.g., ATM, switched Ethernet
- **Internetwork:** set of connected networks
- **Routing:** finding a path through the network between two hosts.

Network Structure

- **Multicast or Broadcast:** send a message to multiple destinations at once.
- **Hardware Broadcast:** use of network-specific features to do broadcast more efficiently than sending to each destination separately.
 - ◆ **Example:** Bus and ring networks support efficient hardware broadcast. Some switched networks and internetworks do, too.
- When the physical network topology is irrelevant, we abstract from it, by including the routing layer in the infrastructure and modeling the network as a complete graph.
 - ◆ **Complete graph:** a graph with an edge between every pair of nodes.

Synchrony (Timing)

- **Terminology** (not standardized): see diagram below
- In contexts when I'm not specifically distinguishing between the application level and the network level, I may use the terms below interchangeably.
- The main hardware elements are labeled in red.



Synchrony

- Main timing-related parameters (in an abstract model), corresponding to main hardware elements:
 - ◆ Process execution speed
 - ◆ Message latency: time between send and deliver
 - ◆ Clock drift rate: ratio of the clock's speed to a perfect clock's speed
- **Asynchronous model:** provides bounds on none of these.
- **Synchronous model:** provides bounds on all of these.
- **Aside:** When referring to **communication primitives** (not timing) synchronous (asynchronous) means the sender does (does not) wait for a reply before proceeding.
- In a **synchronous** system, hosts can maintain **approximately-synchronized clocks**, i.e., there is a bound on the difference between simultaneous values of different clocks.

Is Synchronous or Asynchronous More Realistic?

- Real-time systems and embedded systems are synchronous.
- For general-purpose, multi-user systems, it's hard to get **useful** upper bounds on process execution speed and message latency.
- The upper bounds are much larger than the average/typical values, so algorithms based on waiting until the bound is reached are slow/inefficient.
- Causes of delays of process execution: cache misses, process scheduling, swapping, page faults, disk scheduling
- Causes of delays of message delivery, including message loss and retransmission: collisions in bus networks, contention or buffer overflow in switched networks, processing delays at gateways and re-routing after failures in internetworks

Timed Asynchronous Model

- There exist constants e, l, d such that **usually**

- ◆ Process execution speed is bounded by e
- ◆ Message latency is bounded by l
- ◆ Clock drift is bounded by d

Occasionally some of these bounds are violated.

- Are these “bounds” still useful?
- When the bounds hold, the system is said to be **stable**.
- Design the system to provide stronger guarantees when the system is stable and weaker guarantees when it is not.

Flaviu Cristian and Christof Fetzer. The Timed Asynchronous Distributed System Model. *IEEE Trans. On Parallel and Distributed Systems* 10(6): 642-657 (1999).

Timed Asynchronous Model

● Example: Distributed Database.

- ◆ When the system is **stable**, each query is answered correctly within 1 second.
- ◆ When it is **not stable**, some answers might reflect information that is stale by at most 1 minute.
- ◆ Or: When it is **not stable**, some queries might not be answered. (Client can retry later if desired.)

Failures

- **Benign Failure:** omission failure or timing failure
- **Omission failure:** a component does not perform an expected action at the expected time. (Taxonomy later.)
- **Timing failure**
 - ◆ **Process performance failure:** processes execution speed violates bound
 - ◆ **Channel performance failure:** message latency violates bound
 - ◆ **Clock failure:** clock drift violates bound
- **Arbitrary Failures (a.k.a. Byzantine Failures):** a component behaves arbitrarily, i.e., omits actions and performs actions it should not have performed

Failure Model

- **Failure model** typically specifies
 - ◆ **Types** of failures that may occur
 - ◆ Maximum **number or frequency** of those failures
 - ◆ Maximum **duration** of each failure, if appropriate
- **Transient failure:** failure with finite duration.
 - ◆ Crashes are often considered permanent.
 - ◆ Communication failures are often assumed to be transient.
- A system must satisfy its requirements despite failures allowed by its failure model.

Omission Failures

- **Crash failure:** Faulty process halts forever.
- **Crash and recover**, with stable storage (i.e., storage that survives failures).
 - ◆ Crash and recover without stable storage is equivalent to crash, because a process that recovers without stable storage is equivalent to a new process.
- **Fail-stop:** Faulty process halts forever. Other processes can reliably detect this (no missed failures, no false alarms).

Vassos Hadzilacos and Sam Toueg. Fault-Tolerant Broadcasts and Related Problems. In Sape Mullender, editor, *Distributed Systems*, 2nd edition, chapter 5, pages 97-145. Addison-Wesley, 1993.

Omission Failures

- **Send-omission failure:** process calls send, but message does not make it onto the channel (i.e., is not transmitted onto the network).
- **Link failure:** the communication channel drops messages,
 - ◆ In switched networks, link failure may be due to buffer overflow in a switch, a hardware failure or software bug in a switch, a hardware failure in a cable, etc.
- **Receive-omission failure:** message is received at destination host but not delivered to application, e.g., due to buffer overflow in destination host's network stack.
- Why is it important to distinguish these, even though they have the same effect from the application's point of view?
 - ◆ Different techniques are needed to overcome them.